

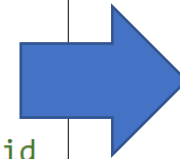
CSE 3100: Systems Programming

Part 2

Lecture 2: More on Processes

Review From Last Lecture (1)

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main()
5  {
6      pid_t value; // process identification used to represent process id
7      value = fork(); //call the fork function to start a separate process
8      printf("In main: value =%d\n", value); //print the id of the process
9  }
```



```
kaleel@CentralCompute:~$ gcc test.c -o test
kaleel@CentralCompute:~$ ./test
In main: value =809
In main: value =0
```

- We discussed the idea of processes and the fork function to create multiple processes.
- When fork is called in the parent code, a child is created which is a clone of the parent. The child starts running AFTER the line where the fork function was called.

Review From Last Lecture (2)

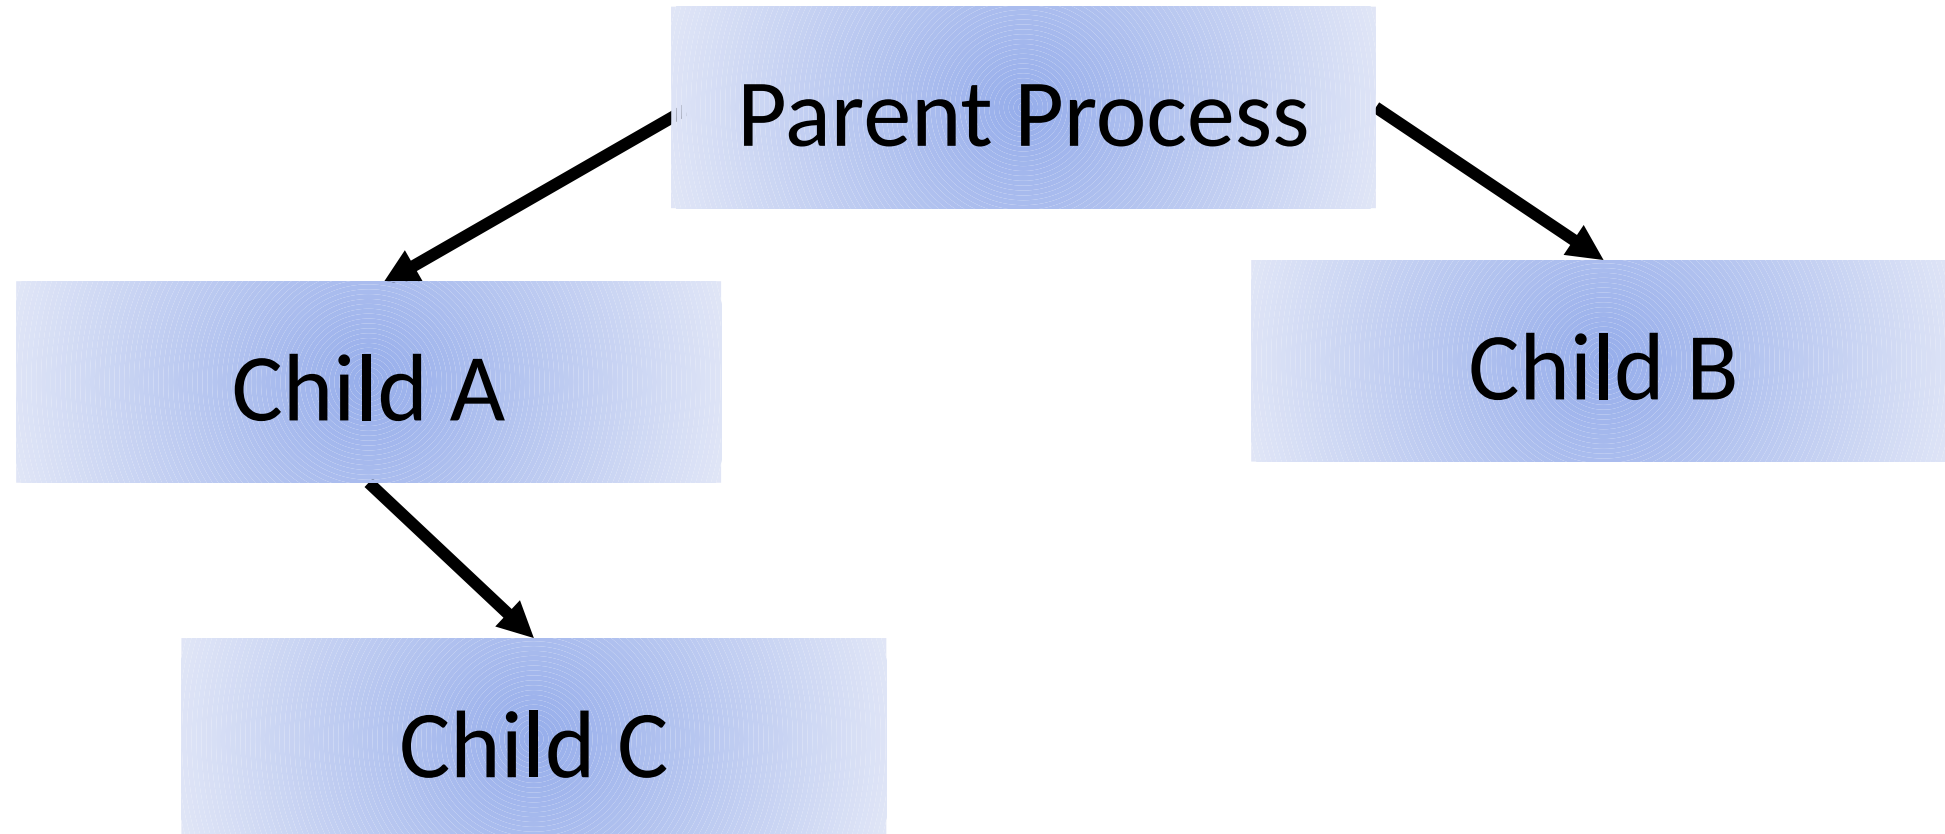
- When we create processes, we usually do so with the goal of making each process run some DIFFERENT piece of code.
- Essentially, we want processes to run in parallel.
- How did we make the parent and child run different tasks? Using the process ID.

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main()
5  {
6      pid_t value; // process identification used to represent process id
7      value = fork(); //call the fork function to start a separate process
8      if(value != 0) //this means we are in the parent process
9      {
10         printf("Let's do the first load of laundry.");
11     }
12     else if(value == 0) //this means we are the child process
13     {
14         printf("Let's do the second load of laundry.");
15     }
16 }
```

Review From Last Lecture (3)

How to keep track of children and parents if you want to call fork multiple times?

- Short Answer: It is not pretty.



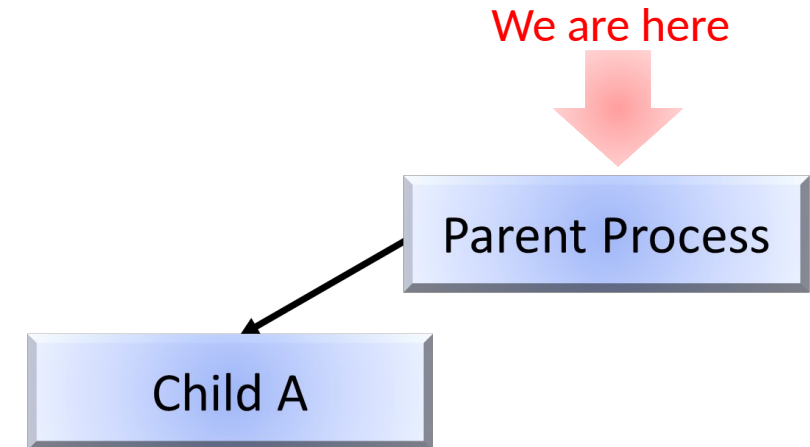
How to keep track of children and parents if you want to call fork multiple times?

- Short Answer: It is not pretty.

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0 ){ //in the parent process
8
9
10
11
12
13
14
15
16     else if(value == 0){
```

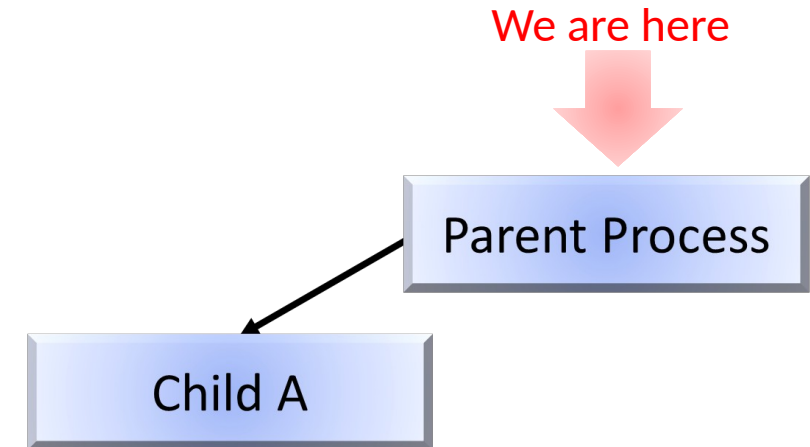
How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
15     }
16     else if(value == 0){
```



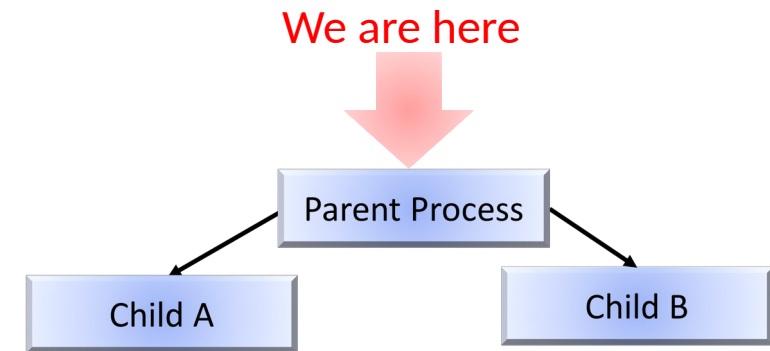
How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
8          value = fork();
9
10
11
12
13
14
15     }
16     else if(value == 0){
```



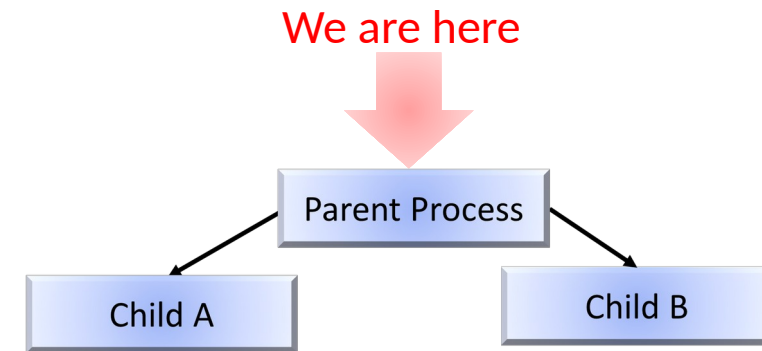
How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
8          value = fork();
9
10
11
12
13
14
15     }
16     else if(value == 0){
```



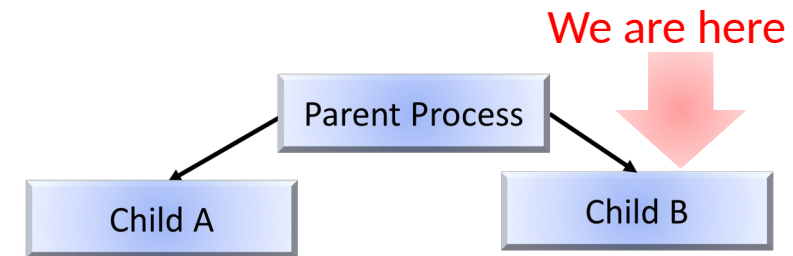
How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
8          value = fork();
9          if(value != 0){ //we are still the parent
10             printf("Parent code finished running here!\n");
11         }
12         else{
13             printf("In the child process B.\n");
14         }
15     }
16     else if(value == 0){
```



How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
8          value = fork();
9          if(value != 0){ //we are still the parent
10             printf("Parent code finished running here!\n");
11         }
12     } else{
13         printf("In the child process B.\n");
14     }
15 }
16 else if(value == 0){
```

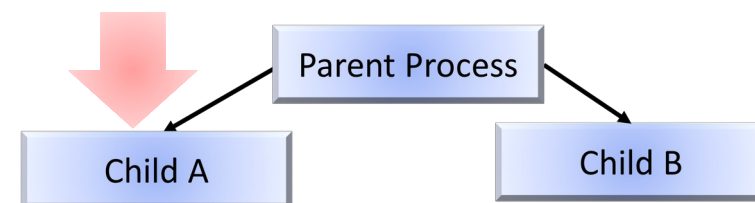


How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0 ){ //in the parent process
```

```
16      else if(value == 0){
```

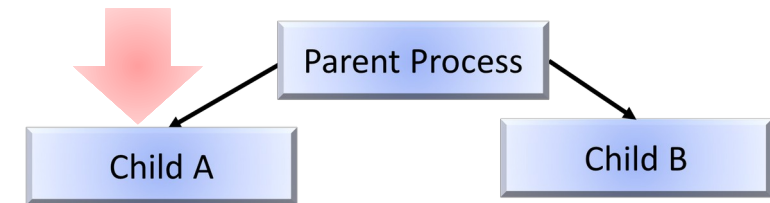
We are here



How to keep track of children and parents if you want to call fork multiple times?

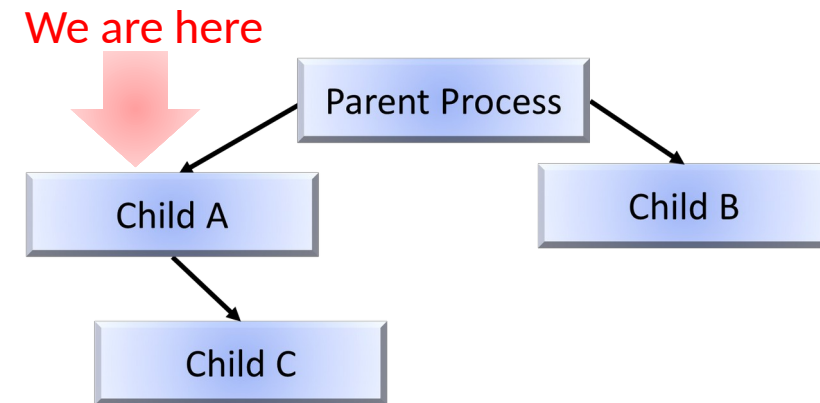
```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
8
9
10
11
12
13
14
15
16     else if(value == 0){
17         printf("In the child process A.\n");
18         value = fork(); //note how the value will get changed in childA when calling fork
```

We are here



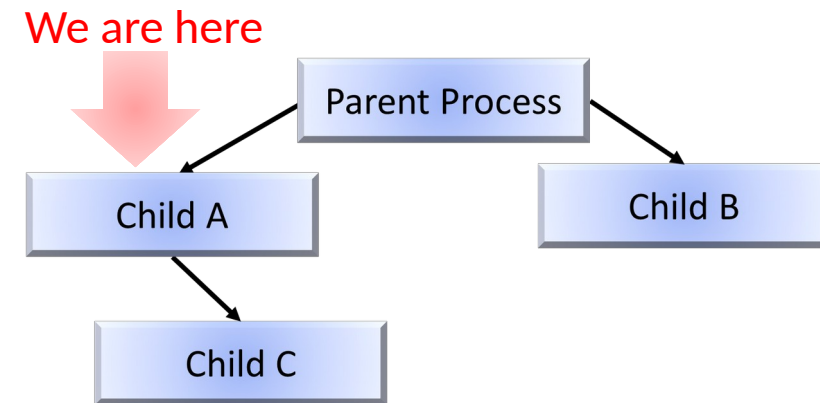
How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0 ){ //in the parent process
8
9
10
11
12
13
14
15
16     else if(value == 0){
17         printf("In the child process A.\n");
18         value = fork(); //note how the value will get changed in childA when calling fork
```



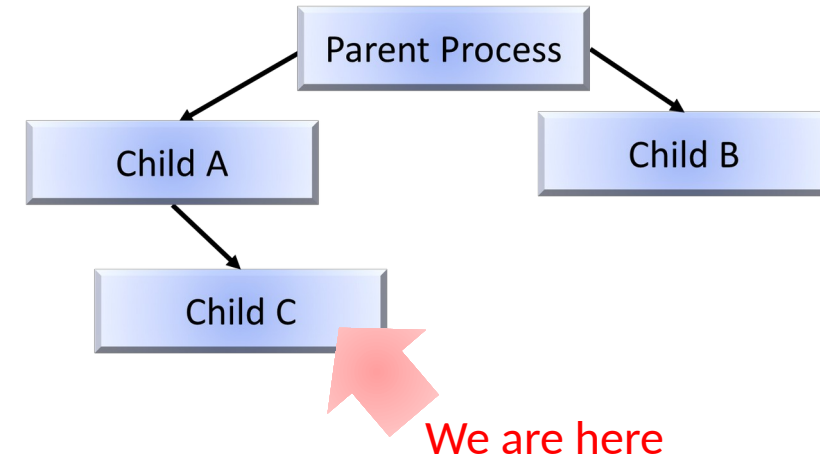
How to keep track of children and parents if you want to call fork multiple times?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0 ){ //in the parent process
8
9
10
11
12
13
14
15
16     else if(value == 0){
17         printf("In the child process A.\n");
18         value = fork(); //note how the value will get changed in childA when calling fork
19         if(value != 0 ){
20             printf("Still inside child process A, but this process is now a parent.\n");
21         }
22     }
23     else if(value == 0){
```



How to keep track of children and parents if you want to call fork multiple times?

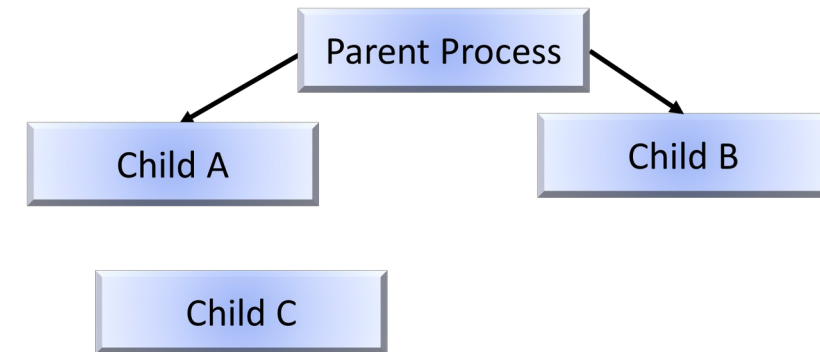
```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0 ){ //in the parent process
8
9
10
11
12
13
14
15
16      else if(value == 0){
17          printf("In the child process A.\n");
18          value = fork(); //note how the value will get changed in childA when calling fork
19          if(value != 0 ){
20              printf("Still inside child process A, but this process is now a parent.\n");
21          }
22          else if(value == 0){
23              printf("Inside child process C.\n");
24          }
25      }
26      return 0;
27  }
```



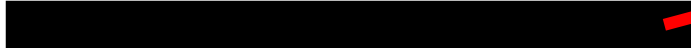
Complete Solution

One more question: could the solution be further improved upon?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(){
5      pid_t value; //process identification used to represent process id
6      value = fork();
7      if(value != 0){ //in the parent process
8          value = fork();
9          if(value != 0){ //we are still the parent
10             printf("Parent code finished running here!\n");
11         }
12         else{
13             printf("In the child process B.\n");
14         }
15     }
16     else if(value == 0){
17         printf("In the child process A.\n");
18         value = fork(); //note how the value will get changed in childA when calling fork
19         if(value != 0){
20             printf("Still inside child process A, but this process is now a parent.\n");
21         }
22         else if(value == 0){
23             printf("Inside child process C.\n");
24         }
25     }
26     return 0;
27 }
```



Question: What if we want the child to run some code that is not contained within the original file?

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     pid_t value; // process identification used to represent process id
7     value = fork(); //call the fork function to start a separate process
8     if(value != 0) //this means we are in the parent process
9     {
10         printf("Let's do the first load of laundry.");
11     }
12     else if(value == 0) //this means we are the child process
13     {
14         
15     }
16 }
```

Run some piece of code that is in a different file.

We'll Explain This in Three Pieces:

First: Let's look at a separate piece of simple code.

Second: How can we run the code without forking?

Third: Run the code WITH forking

First: Let's look at a separate piece of simple code.

adder.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char* argv[]) {
5      int i, sum=0;
6      for(i=1; i<argc; i++)
7          sum += atoi(argv[i]);
8      printf("sum is: %d\n", sum);
9      return 0;
10 }
```

First: Let's look at a separate piece of simple code.

adder.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char* argv[]) {
5      int i, sum=0;
6      for(i=1; i<argc; i++)
7          sum += atoi(argv[i]);
8      printf("sum is: %d\n", sum);
9      return 0;
10 }
```

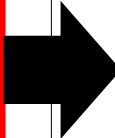
Number of arguments being passed.

- List of numbers to add together, can be any length.
- Numbers are input as strings from the command line.

First: Let's look at a separate piece of simple code.

adder.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char* argv[]) {
5      int i, sum=0;
6      for(i=1; i<argc; i++)
7          sum += atoi(argv[i]);
8      printf("sum is: %d\n", sum);
9      return 0;
10 }
```



- 
- Convert the string numbers to integer data types.
 - Add the numbers to together.
 - Print the result.

Second: How can we run the code without forking?

```
for(i=1;i<argc;i++)  
    sum += atoi(argv[i]);  
printf("sum is: %d\n",sum);
```

```
kaleel@CentralCompute:~$ gcc adder.c -o adder  
kaleel@CentralCompute:~$ ./adder 10 20  
sum is: 30
```

Third: Run the code WITH forking

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./add1";
6      pid_t child = fork();
7      if (child == 0) {
8          
9
10
11
12
13      } else {
14          
15
16
17
18
19      }
20 }
```

Third: Run the code WITH forking

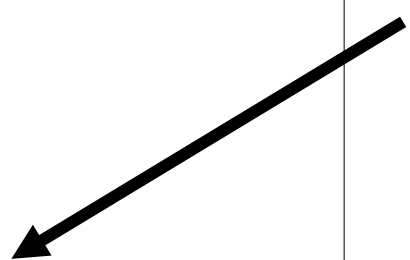
```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./add1";
6      pid_t child = fork();
7      if (child == 0) {
8          [Redacted]
9
10         [Redacted]
11
12         [Redacted]
13     } else {
14         [Redacted]
15
16         [Redacted]
17
18         [Redacted]
19     }
20 }
```

Immediately call the fork function to create two copies of the code.

Third: Run the code WITH forking

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./add1";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1, cmd1, "1", "2", "3", NULL);
10
11
12
13     } else {
14
15
16
17
18
19     }
20 }
```

This is the command
to execute
commands OUTSIDE
of the main code.



Third: Run the code WITH forking

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./add1";
6      pid_t child = fork();
7      if (child == 0) {
8          [REDACTED]
9          [REDACTED]
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         [REDACTED]
15         [REDACTED]
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```

Do some basic
error handling.

What does this look like in practice?



```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1, cmd1, "1", "2", "3", NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1, cmd1, "4", "6", NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```


What does this look like in practice?



```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1, cmd1, "1", "2", "3", NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1, cmd1, "4", "6", NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```


What does this look like in practice?

Parent Process (child variable has non-zero value)



```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```

Child Process (child variable has 0 value)




```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```

What does this look like in practice?

Parent Process (child variable has non-zero value)


```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```



In the parent we will go to line 13
and start executing code.

Child Process (child variable has 0 value)

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```




In the child we will go to line 8
and start executing code.

What does this look like in practice?

Parent Process (child variable has non-zero value)


```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```



At line 15 we'll go into a new code and execute. Its important to note, we DO NOT return.

Child Process (child variable has 0 value)

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```




At line 9 we'll go into a new code and execute. Its important to note, we DO NOT return.

What does this look like in practice?

Parent Process (child variable has non-zero value)

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```




At line 15 we'll go into a new code and execute. Its important to note, we DO NOT return.

We'll get 4+6, which should print 10.

Child Process (child variable has 0 value)

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1,cmd1,"1","2","3",NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1,cmd1,"4","6",NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```

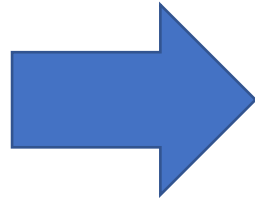


At line 9 we'll go into a new code and execute. Its important to note, we DO NOT return.

We'll get 1+2+3, which should print 6.

What does this look like in practice?

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1, cmd1, "1", "2", "3", NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1, cmd1, "4", "6", NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```



```
kaleel@CentralCompute:~$ ./test
In parent!
In child!
sum is: 10
sum is: 6
```

For this example, where were the files such that they could be

Where were the files such that they could be accessed?

adder.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char* argv[]) {
5      int i, sum=0;
6      for(i=1; i<argc; i++)
7          sum += atoi(argv[i]);
8      printf("sum is: %d\n", sum);
9      return 0;
10 }
```

test.c

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main() {
5      char *cmd1 = "./adder";
6      pid_t child = fork();
7      if (child == 0) {
8          printf("In child!\n");
9          execl(cmd1, cmd1, "1", "2", "3", NULL);
10         printf("Oops.... something went really wrong!\n");
11         perror(cmd1);
12         return -1;
13     } else {
14         printf("In parent!\n");
15         execl(cmd1, cmd1, "4", "6", NULL);
16         printf("Oops.... something went really wrong!\n");
17         perror(cmd1);
18         return -1;
19     }
20 }
```

Main Ubuntu Directory (using Windows Subsystem Linux)

📁 .cache	1/8/2023 10:40 PM	File folder	
📁 .vscode	1/8/2023 10:41 PM	File folder	
📁 .vscode-server	1/8/2023 10:39 PM	File folder	
📄 .bash_history	1/9/2023 7:35 PM	BASH_HISTORY File	1 KB
📄 .bash_logout	1/8/2023 10:38 PM	Bash Logout Sourc...	1 KB
📄 .bashrc	1/8/2023 10:38 PM	Bash RC Source File	4 KB
📄 .motd_shown	1/11/2023 7:17 PM	MOTD_SHOWN File	0 KB
📄 .profile	1/8/2023 10:38 PM	Profile Source File	1 KB
📄 .sudo_as_admin_successful	1/8/2023 10:45 PM	SUDO_AS_ADMIN...	0 KB
📄 .wget-hsts	1/8/2023 10:39 PM	WGET-HSTS File	1 KB
📄 adder	1/11/2023 9:03 PM	File	16 KB
📄 adder	1/11/2023 8:58 PM	C Source File	1 KB
📄 test	1/11/2023 10:33 PM	File	16 KB
📄 test	1/11/2023 10:33 PM	C Source File	2 KB

Process upgrades

- Usually....
 - A fresh clone wants to run *different code*
- This is done by
 - Loading another executable into the process address space
 - [picked up from the file system of course]
- Note
 - **Opened files are NOT AFFECTED** by the upgrade operation



The exec family

Basically any call to Exec1:



- The act of ‘upgrading’ is done by the child with a system call
 - Many variants. "man -S3 execl" for all details

```
#include <unistd.h>
```

```
int execl(const char *path, const char  
*arg0, ... /*, (char *) NULL */ );
```

- **The path** to the executable to load inside our own address space
- A list of arguments to be passed to the new executable
- **A final NULL pointer** to give the “end of argument list”
- If successful, `execl()` **does not return!** Started a new process

How is the executable found?

- Specify a path, like `/bin/ls`
- Specify a file, and the system searches in directories listed in PATH
 - `echo $PATH` in bash to see directories separated by ':'

```
int execl(const char *path, const char *arg0, ...  
          /*, (char *) NULL */ );
```

`// execlp() searches paths for file`

```
int execlp(const char *file, const char *arg0, ...  
          /*, (char *) NULL */ );
```

The `exec` family 2

// If the number of arguments is unknown at compile time

```
#include <unistd.h>
```

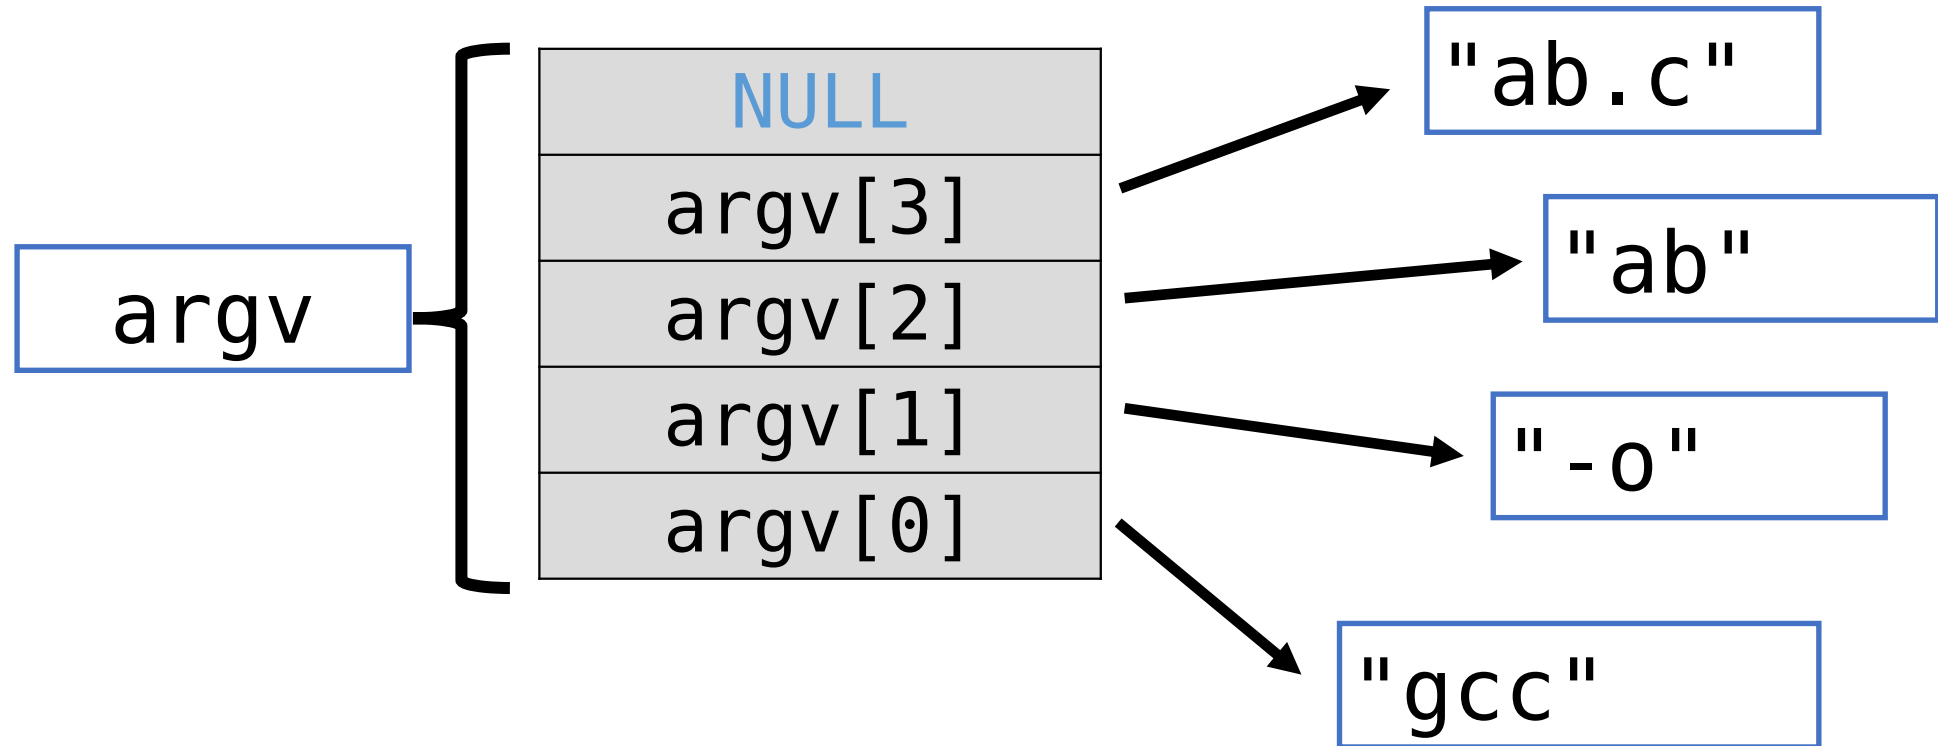
```
int execlv(const char *path, char *const argv[]);
```

```
int execlvp(const char *file, char *const  
argv[]);
```

- The arguments in `execl` in are placed in an array
 - `argv` is the `argv` you see in the main function!
- `execlv` needs a path while `execlvp` can search file in `PATH`
- Start a new process if successful. Similar to `execl`

A small note on: argv to execv and execvp

```
$gcc -o ab ab.c
```



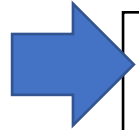
- Note that the last value in `argv` is null? Why?
- So that the code knows where to stop reading.

Common Errors in Writing Forking Code (Example 1)

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork()
fails
} else if (pid == 0) {
    child_tasks();
} else {
    parent_tasks();
}
more_parent_tasks();
```

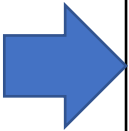
What is wrong with this code?

Example 1: Consider the child process



```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork()
fails
} else if (pid == 0) {
    child_tasks();
} else {
    parent_tasks();
}
more_parent_tasks();
```


Example 1: Consider the child process



```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork()
fails
} else if (pid == 0) {
    child_tasks();
} else {
    parent_tasks();
}
more_parent_tasks();
```

Example 1: Consider the child process

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork()
fails
} else if (pid == 0) {
    child_tasks();
} else {
    parent_tasks();
}
more_parent_tasks();
```



Example 1: Consider the child process

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork()
fails
} else if (pid == 0) {
    child_tasks();
} else {
    parent_tasks();
}
more_parent_tasks();
```

Visualization of the Child Process
when you ask it to mow the lawn:



Example 1: Corrected Code

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork() fails
} else if (pid == 0) {
    child_tasks();
    exit(0);    // terminate the child process
} else {
    parent_tasks();
}
more_parent_tasks();
```

Example 2: What could go wrong?

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork() fails
}
else if (pid == 0) {
    // in child process
    execlp("genie", "genie", "clean the house", NULL);
}
// in parent process
online_shopping();
```

*Would you ever know if
execlp failed?*

Example 2: Corrected Code

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1);    // exit if fork() fails
}
else if (pid == 0) {
    // in child process
    execlp("genie", "genie", "clean the house", NULL);
    printf("Something went really wrong in child process!\n");
}
// in parent process
online_shopping();
```

File APIs

(API = Application Programming Interface)

- Remember the (C standard library) IO APIs
 - The “f” family (fopen, fclose, fread, fgetc, fscanf, fprintf,...)
 - All these use a FILE* abstraction to represent a file
 - Additional features: user-space buffering, line-ending translation, formatted I/O, etc.
- UNIX has lower-level APIs for file handling
 - Directly mapped to **system calls**
 - open, close, read, ...
 - Use *file descriptors* [which are just **integers**]
 - Deal with bytes only

Some low level file APIs

- Read the man pages (man -s2 ...) for more functions

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

int open(const char *path, int oflag);
int open(const char *path, int oflag, int mode);
int close(int fd);

ssize_t read(int fd, void *buf, size_t nbyte);
ssize_t write(int fd, const void *buf, size_t nbyte);

off_t lseek(int fd, off_t offset, int whence);
```

How to open a file

```
#include <fcntl.h>  
#include <unistd.h>
```

```
int open(const char *path, int oflag);
```

- Parameters
 - path: the path to the file to be opened/created
 - oflag: read, write, or read and write, and more (on the next slide)
- The function returns [a file descriptor](#), a small, nonnegative integer
 - Return -1 on error

Flags in open()

- Must include one of the following:
`O_RDONLY` (read only), `O_WRONLY` (write only), or `O_RDWR` (read and write)
- And or-ed (|) with many optional flags, for example,
 - `O_TRUNC`: Truncate the file (remove existing contents) if opening a file for write
 - `O_CREAT`: Create a file if it does not exist.

Example:

```
// remember open() returns -1 on error
fd1 = open("a.txt", O_RDONLY); // open for read
fd1 = open("a.txt", O_RDWR); // open for read and write
fd1 = open("a.txt", O_RDWR|O_TRUNC); // read, write, truncate the
file
```

Creating a file with open()

```
// a mode must be provided if O_CREAT or O_TMPFILE is set
```

```
int open(const char *path, int oflag, int mode);
```

mode: specify permissions when a new, or temporary, file is created.

```
open("b.txt", O_WRONLY|O_TRUNC|O_CREAT, 0600);
```

```
// open b.txt for write. If the file exists, clear (truncate) the contents.
```

```
// if the file does not exist, create one, and set the permission so that the owner of the file can read and write, but other people cannot.
```

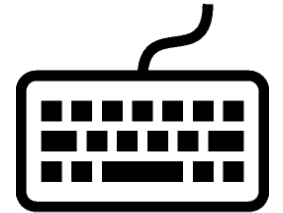
File descriptors

- A file descriptor is a nonnegative integer associated with a file.

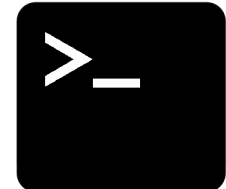
```
#include <stdio.h>
int fileno(FILE *stream);
// returns a file descriptor for a stream
```

FD	FILE *
0	stdin
1	stdout
2	stderr

The input, for example the process takes inputs from the keyboard.



Where the process gives outputs e.g. the terminal.



Where the process reports errors, e.g. could be a log file recording the errors.



File descriptors after fork and exec

- Opened files are NOT AFFECTED by the upgrade operation

```
pid_t pid = fork();
assert(pid >= 0);
if (pid == 0) {
    // Child process can access FDs 0, 1, and 2
    // if execl() is successful, gcc can access FDs 0, 1, and 2
    execlp("gcc", "gcc", "a.c", NULL);
    // If control gets here, execlp() failed.
    // Remember to terminate the child process!
    return 1;
}
```

Visualization of File Descriptor with Fork

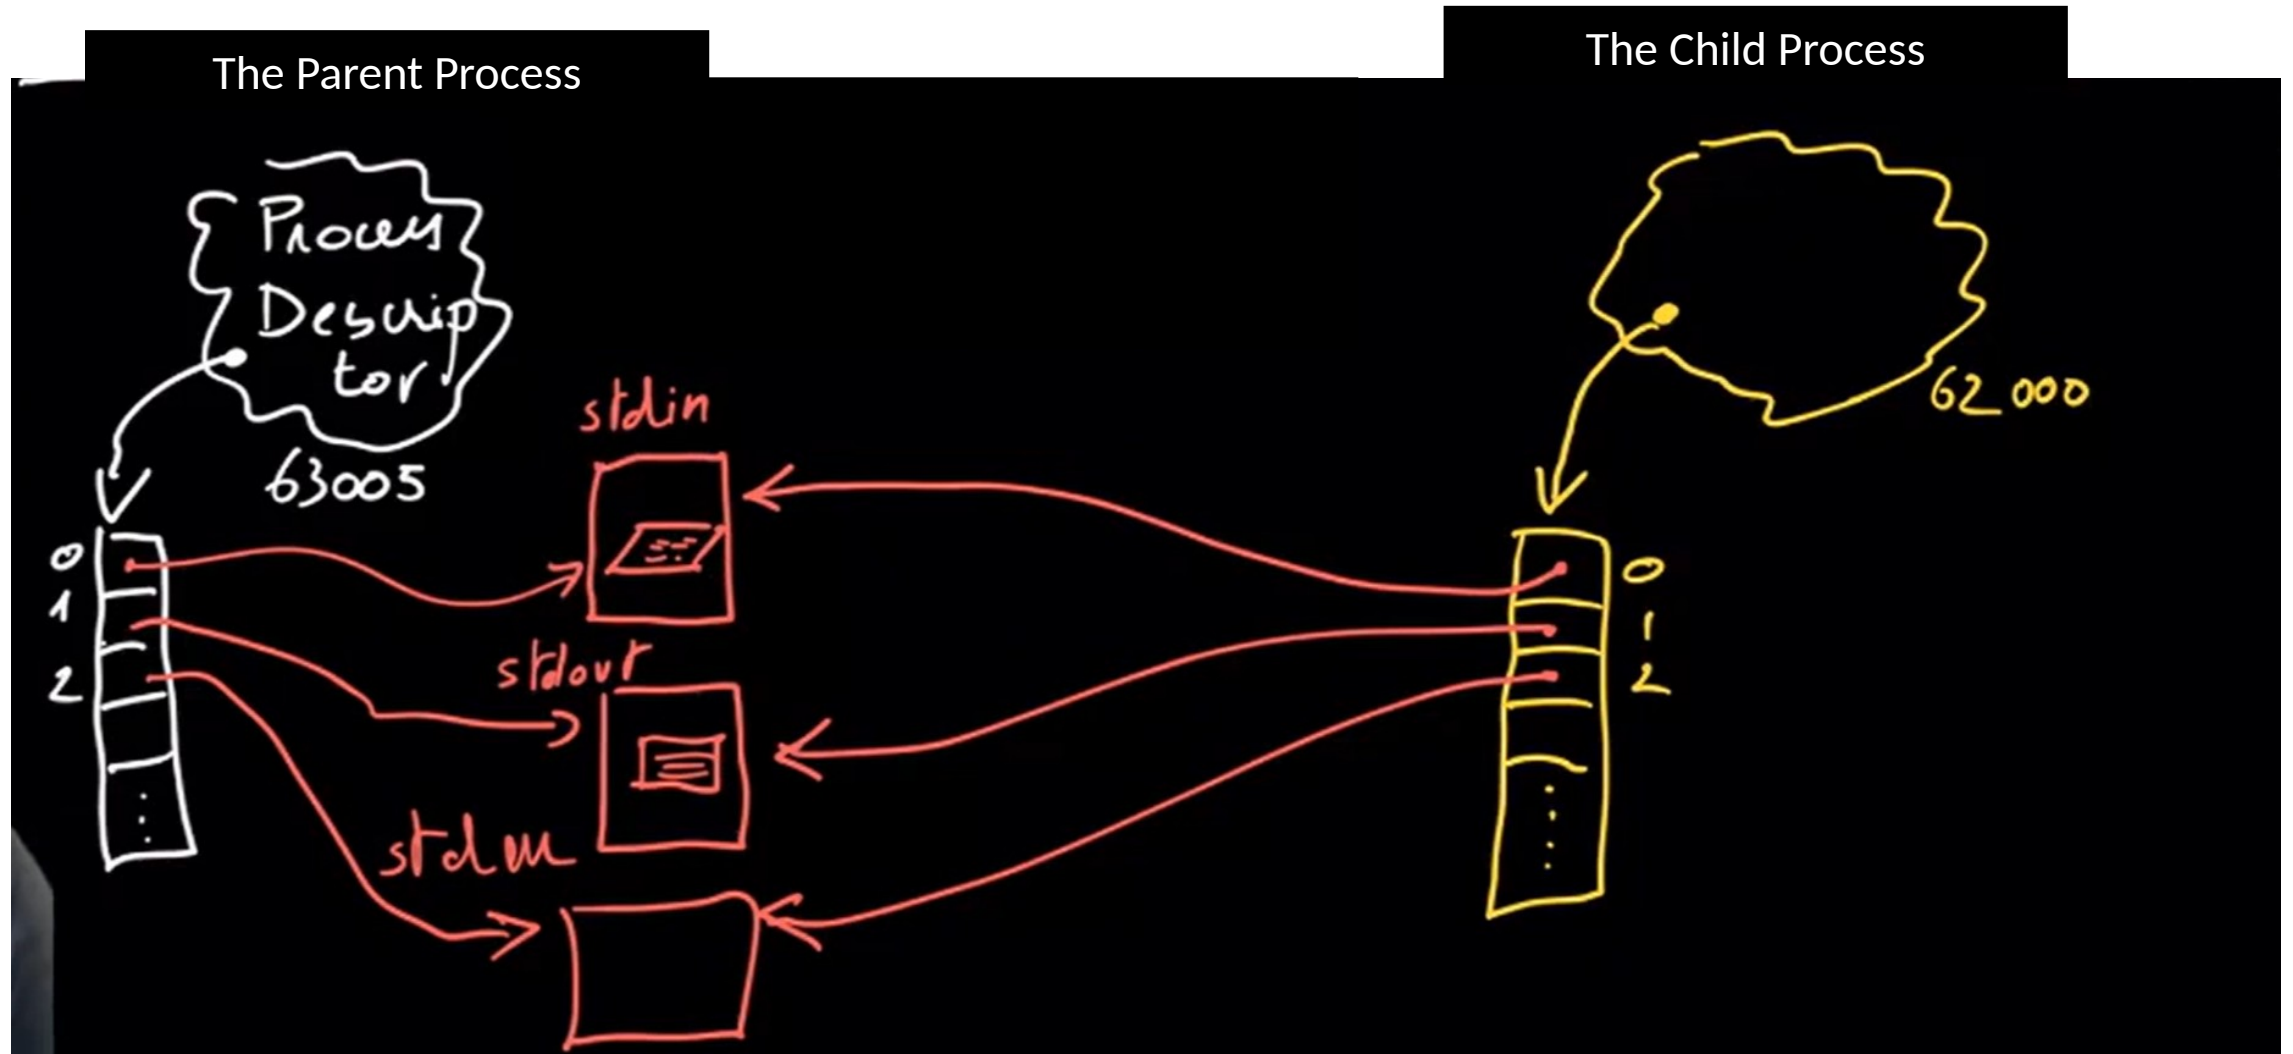
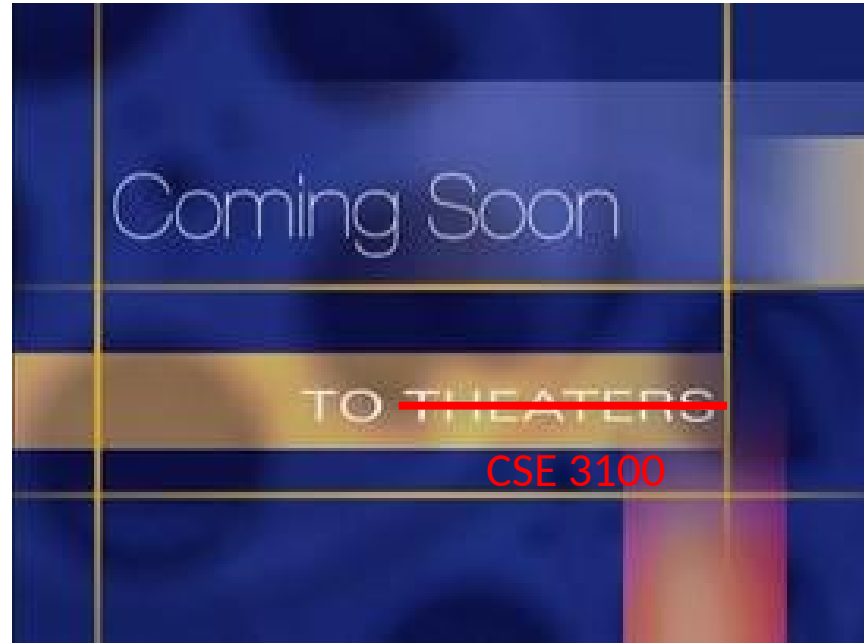


Figure Source: Professor Laurent Michel Youtube Lecture Videos

Preview of the future



- You already know that child processes copy the variables from the parent code.
- It follows that the file descriptors are also copied over to the child.
- So why all the focus on file descriptors? This will come in handy in the future when we want processes to communicate with each other.

Figure Sources

1. <https://i.kym-cdn.com/entries/icons/facebook/000/019/404/upgraddddd.jpg>
2. <https://www.mememaker.net/static/images/memes/4761591.jpg>
3. <https://thumbs.gfycat.com/HollowPositiveAnophelesmosquito-max-1mb.gif>
4. <https://iconarchive.com/download/i87838/icons8/ios7/Computer-Hardware-Keyboard.ico>
5. <https://icons.iconarchive.com/icons/paomedia/small-n-flat/1024/terminal-icon.png>
6. <https://cdn-icons-png.flaticon.com/512/1388/1388902.png>
7. <https://static.seekingalpha.com/uploads/2013/3/21/7360901-13638972437431467-Robert-Wagner.jpg>